

Exercise from chap01ex.ipynb, exclude the prefilled book demo code

```
from os.path import basename, exists

def download(url):
    filename = basename(url)
    if exists(filename):
        from urllib.request import urlretrieve

        local, _ = urlretrieve(url, filename)
        print("Downloaded " + local)

download("https://github.com/AllenDowney/ThinkStats2/raw/master/code/
thinkstats2.py")
download("https://github.com/AllenDowney/ThinkStats2/raw/master/code/
thinkplot.py")
download("https://github.com/AllenDowney/ThinkStats2/raw/master/code/
nsfg.py")

Downloaded thinkstats2.py
Downloaded thinkplot.py
Downloaded nsfg.py

download("https://github.com/AllenDowney/ThinkStats2/raw/master/code/
2002FemPreg.dct")
download(

"https://github.com/AllenDowney/ThinkStats2/raw/master/code/2002FemPre
g.dat.gz"
)

Downloaded 2002FemPreg.dct
Downloaded 2002FemPreg.dat.gz
```

Read NSFG data into a Pandas DataFrame.

```
import nsfg

preg = nsfg.ReadFemPreg()

C:\Users\gyanr\gyan-python-workspace\DSC-530\nsfg.py:68:
FutureWarning: A value is trying to be set on a copy of a DataFrame or
```

Series through chained assignment using an inplace method. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df.birthwgt_lb.replace(na_vals, np.nan, inplace=True)
```

C:\Users\gyanr\gyan-python-workspace\DSC-530\nsfg.py:69:
FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df.birthwgt_oz.replace(na_vals, np.nan, inplace=True)
```

C:\Users\gyanr\gyan-python-workspace\DSC-530\nsfg.py:70:
FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df.hpagelb.replace(na_vals, np.nan, inplace=True)
```

C:\Users\gyanr\gyan-python-workspace\DSC-530\nsfg.py:72:
FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the

original object.

```
df.babysex.replace([7, 9], np.nan, inplace=True)
```

C:\Users\gyanr\gyan-python-workspace\DSC-530\nsfg.py:73:
FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df.nbrnaliv.replace([9], np.nan, inplace=True)
```

exercise 1.1 start from here:

Select the `birthord` column, print the value counts, and compare to results published in the [codebook](#)

```
birthord = preg['birthord']
preg['totalwgt_lb'].describe()
count      9038.000000
mean       7.265628
std        1.408293
min         0.125000
25%        6.500000
50%        7.375000
75%        8.125000
max        15.437500
Name: totalwgt_lb, dtype: float64
```

We can also use `isnull` to count the number of nans.

```
preg.birthord.isnull().sum()
4445
```

Select the `prglngth` column, print the value counts, and compare to results published in the [codebook](#)

```
preg.prglngth.value_counts().sort_index()
```

prglngh

0	15
1	9
2	78
3	151
4	412
5	181
6	543
7	175
8	409
9	594
10	137
11	202
12	170
13	446
14	29
15	39
16	44
17	253
18	17
19	34
20	18
21	37
22	147
23	12
24	31
25	15
26	117
27	8
28	38
29	23
30	198
31	29
32	122
33	50
34	60
35	357
36	329
37	457
38	609
39	4744
40	1120
41	591
42	328
43	148
44	46
45	10
46	1
47	1
48	7

```
50      2
Name: count, dtype: int64
```

To compute the mean of a column, you can invoke the `mean` method on a Series. For example, here is the mean birthweight in pounds:

```
preg.totalwgt_lb.mean()
7.265628457623368
```

Create a new column named `totalwgt_kg` that contains birth weight in kilograms. Compute its mean. Remember that when you create a new column, you have to use dictionary syntax, not dot notation.

```
preg['totalwgt_kg'] = preg.totalwgt_lb*0.45359237
preg.totalwgt_kg.mean()
3.295633631632828
```

`nsfg.py` also provides `ReadFemResp`, which reads the female respondents file and returns a `DataFrame`:

```
download("https://github.com/AllenDowney/ThinkStats2/raw/master/code/
2002FemResp.dct")
download("https://github.com/AllenDowney/ThinkStats2/raw/master/code/
2002FemResp.dat.gz")

Downloaded 2002FemResp.dct
Downloaded 2002FemResp.dat.gz

resp = nsfg.ReadFemResp()
```

`DataFrame` provides a method `head` that displays the first five rows:

```
resp.head()
   caseid  rscrinf  rdormres  rostscrn  rscreenhisp  rscreenrace
age_a \
0      2298      1        5        5          1          5.0
27
1      5012      1        5        1          5          5.0
42
2     11586      1        5        1          5          5.0
43
3      6794      5        5        4          1          5.0
15
4       616      1        5        4          1          5.0
20
```

	age_r	cmbirth	agescrn	...	pubassis_i	basewgt
adj_mod_basewgt \						
0	27	902	27	...	0	3247.916977
5123.759559						
1	42	718	42	...	0	2335.279149
2846.799490						
2	43	708	43	...	0	2335.279149
2846.799490						
3	15	1042	15	...	0	3783.152221
5071.464231						
4	20	991	20	...	0	5341.329968
6437.335772						

	finalwgt	secu_r	sest	cmintvw	cmlstyr	screen time	intvlngth
0	5556.717241	2	18	1234	1222	18:26:36	110.492667
1	4744.191350	2	18	1233	1221	16:30:59	64.294000
2	4744.191350	2	18	1234	1222	18:19:09	75.149167
3	5923.977368	2	18	1234	1222	15:54:43	28.642833
4	7229.128072	2	18	1233	1221	14:19:44	69.502667

[5 rows x 3087 columns]

Select the `age_r` column from `resp` and print the value counts. How old are the youngest and oldest respondents?

```
resp.age_r.value_counts().sort_index()
```

age_r	
15	217
16	223
17	234
18	235
19	241
20	258
21	267
22	287
23	282
24	269
25	267
26	260
27	255
28	252
29	262
30	292

```

31    278
32    273
33    257
34    255
35    262
36    266
37    271
38    256
39    215
40    256
41    250
42    215
43    253
44    235
Name: count, dtype: int64

```

We can use the `caseid` to match up rows from `resp` and `preg`. For example, we can select the row from `resp` for `caseid` 2298 like this:

```

resp[resp.caseid==2298]

```

	caseid	rscrinf	rdormres	rostscrn	rscreenhisp	rscreenrace
age_a \						
0	2298	1	5	5	1	5.0
27						

	age_r	cmbirth	agescrn	...	pubassis_i	basewgt
adj_mod_basewgt \						
0	27	902	27	...	0	3247.916977
5123.759559						

	finalwgt	secu_r	sest	cmintvw	cmlstyr	screentime	intvlngth
0	5556.717241	2	18	1234	1222	18:26:36	110.492667


```

[1 rows x 3087 columns]

```

And we can get the corresponding rows from `preg` like this:

```

preg[preg.caseid==2298]

```

	caseid	pregordr	howpreg_n	howpreg_p	moscurrp	nowprgdk
pregendl \						
2610	2298	1	NaN	NaN	NaN	NaN
6.0						
2611	2298	2	NaN	NaN	NaN	NaN
6.0						
2612	2298	3	NaN	NaN	NaN	NaN

```

6.0
2613      2298          4      NaN      NaN      NaN      NaN
6.0

      pregend2  nbrnaliv  multbrth  ...  religion_i  metro_i
basewgt \
2610      NaN      1.0      NaN  ...      0      0
3247.916977
2611      NaN      1.0      NaN  ...      0      0
3247.916977
2612      NaN      1.0      NaN  ...      0      0
3247.916977
2613      NaN      1.0      NaN  ...      0      0
3247.916977

      adj_mod_basewgt      finalwgt  secu_p  sest  cmintvw  totalwgt_lb
\
2610      5123.759559  5556.717241      2    18      NaN      6.8750
2611      5123.759559  5556.717241      2    18      NaN      5.5000
2612      5123.759559  5556.717241      2    18      NaN      4.1875
2613      5123.759559  5556.717241      2    18      NaN      6.8750

      totalwgt_kg
2610      3.118448
2611      2.494758
2612      1.899418
2613      3.118448

[4 rows x 245 columns]

```

How old is the respondent with `caseid` 1?

```

resp[resp.caseid==1].age_r
1069      44
Name: age_r, dtype: int64

```

What are the pregnancy lengths for the respondent with `caseid` 2298?

```

preg[preg.caseid==2298].prglngth
2610      40
2611      36
2612      30
2613      40
Name: prglngth, dtype: int64

```


What was the birthweight of the first baby born to the respondent with `caseid` 5012?

```
preg[preg.caseid==5012]
```

```
      caseid  pregordr  howpreg_n  howpreg_p  moscurrp  nowprgdk
pregend1 \
5515    5012         1         NaN         NaN         NaN         NaN
6.0
```

```
      pregend2  nbrnaliv  multbrth  ...  religion_i  metro_i
basewgt \
5515         NaN        1.0        NaN  ...         0         0
2335.279149
```

```
      adj_mod_basewgt  finalwgt  secu_p  sest  cmintvw  totalwgt_lb
\
5515         2846.79949  4744.19135         2    18         NaN         6.0
```

```
      totalwgt_kg
5515         2.721554
```

```
[1 rows x 245 columns]
```

```
preg[(preg.caseid==5012) & (preg.pregordr==1)].totalwgt_lb
```

```
5515    6.0
Name: totalwgt_lb, dtype: float64
```

```
preg.describe()
```

```
      caseid  pregordr  howpreg_n  howpreg_p
moscurrp \
count  13593.000000  13593.000000  352.000000  349.000000  352.000000
mean    6216.526595    2.349150  15.144886    1.34384    4.647727
std     3645.417341    1.577807  13.922211    0.47567    2.527523
min       1.000000    1.000000    0.000000    1.00000    0.000000
25%     3022.000000    1.000000    5.000000    1.00000    2.000000
50%     6161.000000    2.000000    9.000000    1.00000    5.000000
75%     9423.000000    3.000000   23.000000    2.00000    7.000000
max    12571.000000   19.000000   99.000000    2.00000    9.000000
```

```
      nowprgdk  pregend1  pregend2  nbrnaliv  multbrth  ...
\
```

count	3.000000	13241.000000	18.000000	9144.000000	163.000000	...
mean	3.666667	4.650177	4.055556	1.022419	1.834356	...
std	4.618802	1.849790	1.696787	0.190098	1.630208	...
min	1.000000	1.000000	1.000000	1.000000	1.000000	...
25%	1.000000	3.000000	3.000000	1.000000	1.000000	...
50%	1.000000	6.000000	4.000000	1.000000	1.000000	...
75%	5.000000	6.000000	6.000000	1.000000	1.000000	...
max	9.000000	9.000000	6.000000	5.000000	5.000000	...

	religion_i	metro_i	basewgt	adj_mod_basewgt
finalwgt \				
count	13593.000000	13593.0	13593.000000	13593.000000
13593.000000				
mean	0.003016	0.0	4216.271164	5383.982581
8196.422280				
std	0.058727	0.0	3982.680473	5640.499431
9325.918114				
min	0.000000	0.0	64.577101	71.201194
118.656790				
25%	0.000000	0.0	2335.445237	2798.048902
3841.375308				
50%	0.000000	0.0	3409.648504	4127.220642
6256.592133				
75%	0.000000	0.0	4869.941451	5795.692880
9432.360931				
max	2.000000	0.0	99707.832014	157143.686687
261879.953864				

	secu_p	sest	cmintvw	totalwgt_lb	totalwgt_kg
count	13593.000000	13593.000000	0.0	9038.000000	9038.000000
mean	1.487310	44.083352	NaN	7.265628	3.295634
std	0.499857	24.110403	NaN	1.408293	0.638791
min	1.000000	1.000000	NaN	0.125000	0.056699
25%	1.000000	25.000000	NaN	6.500000	2.948350
50%	1.000000	45.000000	NaN	7.375000	3.345244
75%	2.000000	65.000000	NaN	8.125000	3.685438
max	2.000000	84.000000	NaN	15.437500	7.002332

[8 rows x 245 columns]

End of exercise 1.1

Exercise 1.2: Create a file named chap01ex.py and write code that reads the respondent file, 2002FemResp.dat.gz

```
# Value counts for 'pregnum'

import nsfg

# Load the datasets
resp = nsfg.ReadFemResp()
preg = nsfg.ReadFemPreg()

# Print the value counts for 'pregnum' variable in the 'resp' dataset
print("Value counts for 'pregnum' in the 'resp' dataset:")
print(resp['pregnum'].value_counts())

# Create a dictionary mapping caseid to a list of indices into the
# pregnancy DataFrame
preg_map = nsfg.MakePregMap(preg)

# Create a new column in 'resp' dataset representing the number of
# records in the pregnancy DataFrame for each caseid
resp['preg_count'] = resp['caseid'].map(lambda x: len(preg_map.get(x,
[])))

# Compare 'pregnum' for each respondent with the number of records in
# the pregnancy
comparison_result = resp['pregnum'] == resp['preg_count']

# Print the result
print("\nComparison result:")
print(comparison_result.value_counts())
```

Value counts for 'pregnum' in the 'resp' dataset:

pregnum	
0	2610
2	1432
1	1267
3	1110
4	611
5	305
6	150
7	80
8	40

```
9      21
10     9
11     3
12     2
14     2
19     1
Name: count, dtype: int64
```

```
Comparison result:
True      7643
Name: count, dtype: int64
```

C:\Users\gyanr\gyan-python-workspace\DSC-530\nsfg.py:68:
FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df.birthwgt_lb.replace(na_vals, np.nan, inplace=True)
C:\Users\gyanr\gyan-python-workspace\DSC-530\nsfg.py:69:
FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df.birthwgt_oz.replace(na_vals, np.nan, inplace=True)
C:\Users\gyanr\gyan-python-workspace\DSC-530\nsfg.py:70:
FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the

original object.

```
df.hpagelb.replace(na_vals, np.nan, inplace=True)
```

C:\Users\gyanr\gyan-python-workspace\DSC-530\nsfg.py:72:
FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df.babysex.replace([7, 9], np.nan, inplace=True)
```

C:\Users\gyanr\gyan-python-workspace\DSC-530\nsfg.py:73:
FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df.nbrnaliv.replace([9], np.nan, inplace=True)
```

Summary:

For most respondents, the recorded number of pregnancies in the 'resp' dataset aligns with the actual number of pregnancy records in the pregnancy dataset. The absence of False values in the comparison result suggests that there are no inconsistencies between the reported number of pregnancies and the actual number of pregnancy records for any respondent. In summary, based on the comparison result, the 'pregnum' variable in the 'resp' dataset appears to be consistent with the number of pregnancy records in the pregnancy dataset, indicating a high level of data integrity and reliability in the reported pregnancy information.

End of exercise 1.2

```
from scipy.stats import ttest_ind

# Extracting the 'totalwgt_lb' column and 'pregordr' column
totalwgt_lb = preg.totalwgt_lb
pregordr = preg.pregordr
```

```

# Separating first babies and others
first_babies = totalwgt_lb[pregrordr == 1]
others = totalwgt_lb[pregrordr > 1]

# Calculating means and standard deviations for both groups
mean_first_babies = first_babies.mean()
mean_others = others.mean()
std_first_babies = first_babies.std()
std_others = others.std()

# Independent two-sample t-test
t_statistic, p_value = ttest_ind(first_babies, others,
equal_var=False)

# Cohen's d
cohens_d = (mean_first_babies - mean_others) / (((std_first_babies**2
+ std_others**2) / 2)**0.5)

print("Mean weight of first babies:", mean_first_babies)
print("Mean weight of others:", mean_others)
print("Cohen's d:", cohens_d)
print("p-value:", p_value)

print("Mean weight of first babies:", mean_first_babies)
print("Mean weight of others:", mean_others)

Mean weight of first babies: 7.204107733975324
Mean weight of others: 7.301399825021872
Cohen's d: -0.06904986139204121
p-value: nan
Mean weight of first babies: 7.204107733975324
Mean weight of others: 7.301399825021872

```

2-1 (Based on the results in this chapter, suppose you were asked to summarize what you learned about whether first babies arrive late...)

Mean pregnancy length for first babies is 38.601;

for other babies it is 38.523.

the standard deviation is 2.7 weeks

(Which summary statistics would you use if you wanted to get a story on the evening news? Which ones would you use if you wanted to reassure an anxious patient?)

Although standard deviation makes sense, it would be appropriate to highlight key summary statistics that capture the overall distribution of baby weights. The mean (average) weight of all babies is 7.27 pounds, with a standard deviation of 1.41 pounds. The range of weights spans from 0.13 to 15.44 pounds. This information provides a general overview of baby weights and could be emphasized to give viewers a sense of the typical weight range and variation.

End of exercise 2-1

2-4 (Using the variable `totalwgt_lb`, investigate whether first babies are lighter or heavier than others...)

```
import math
import matplotlib.pyplot as plt
import nsfg

def read_data():
    preg = nsfg.ReadFemPreg()

    live = preg[preg.outcome == 1]
    firsts = live[live.birthord == 1]
    others = live[live.birthord != 1]

    return live, firsts, others

def calculate(live, firsts, others):

    mean1 = firsts.totalwgt_lb.mean()
    mean2 = others.totalwgt_lb.mean()
    print(f'Mean of First {mean1=} , others {mean2=}')

    var1 = firsts.totalwgt_lb.var()
    var2 = others.totalwgt_lb.var()
    print(f'Variance of First {var1=} , others {var2=}')

    # calculate Choen D
    diff = mean1 - mean2
    n1, n2 = len(firsts.totalwgt_lb), len(others.totalwgt_lb)
```

```

pooled_var = (n1 * var1 + n2 * var2) / (n1 + n2)
d = diff / math.sqrt(pooled_var)
print('Cohen d', d)

plt.hist(firsts.totalwgt_lb, label='first baby', alpha=0.5)
plt.hist(others.totalwgt_lb, label='other babies', alpha=0.5)
plt.legend(loc='best')
plt.show()

```

```

def main(script):
    live, firsts, others = read_data()
    calculate(live, firsts, others)

```

```

if __name__ == '__main__':
    main('test')

```

Mean of First mean1=7.201094430437772 , others 7.325855614973262
 Variance of First var1=2.018027300915786 , others 1.9437810258964716
 Cohen d -0.08864367587767717

C:\Users\gyanr\gyan-python-workspace\DSC-530\nsfg.py:68:
 FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
 The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```

df.birthwgt_lb.replace(na_vals, np.nan, inplace=True)

```

C:\Users\gyanr\gyan-python-workspace\DSC-530\nsfg.py:69:
 FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
 The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```

df.birthwgt_oz.replace(na_vals, np.nan, inplace=True)

```

C:\Users\gyanr\gyan-python-workspace\DSC-530\nsfg.py:70:
 FutureWarning: A value is trying to be set on a copy of a DataFrame or

Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df.hpagelb.replace(na_vals, np.nan, inplace=True)
```

C:\Users\gyanr\gyan-python-workspace\DSC-530\nsfg.py:72:
FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.


For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df.babysex.replace([7, 9], np.nan, inplace=True)
```

C:\Users\gyanr\gyan-python-workspace\DSC-530\nsfg.py:73:
FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df.nbrnaliv.replace([9], np.nan, inplace=True)
```



Mean

First babies 7.201094430437772

Others 7.325855614973262

Variance

First babies 2.0180273009157768

Others 1.9437810258964572

Difference in lbs -0.12476118453549034

Difference in oz -1.9961789525678455

Difference relative to mean (%age points) -1.7171423678372415

Cohen d -0.088672927072602

first babies is slightly lighter (~2 oz) than others, but the difference is just a fraction of the first babies mean (less than 2%). However, Based on the Cohen-D value, the variation on baby weight is bigger than the variation in pregnancy length