

```

from __future__ import print_function, division
import numpy as np
import pandas as pd
import seaborn as sns
from matplotlib import pyplot as plt
import thinkstats2
import thinkplot
import hypothesis
import regression
import statsmodels.formula.api as smf
import warnings
warnings.filterwarnings("ignore")

all_health_df = pd.read_csv('datasets/h2_ins.csv')
all_health_df = all_health_df.drop(columns=['index'], errors='ignore')

all_health_df.shape
all_health_df.head(3)

```

	age	sex	bmi	children	smoker	charges	city
region							
0	19	female	27.90	0	yes	16884.9240	Beechwood
1	18	male	33.77	1	no	1725.5523	Louisville
2	28	male	33.00	3	no	4449.4620	Merriam

A minimum of 5 variables in your dataset used during your analysis (for help with selecting, the author made his selection on page 6 of your book). Consider what you think could have an impact on your question – remember this is never perfect, so don't be worried if you miss one (Chapter 1).

```

all_health_df.columns

Index(['age', 'sex', 'bmi', 'children', 'smoker', 'charges', 'city',
       'region'], dtype='object')

```

Describe what the 5 variables mean in the dataset (Chapter 1).

Age- Age of the insured

Sex- The gender of the insured

Bmi- The body mass index for the insured

Children - Number of children of the insured.

Smoker - Whether the insured smokes or not

Charges - What is the medical expenses claimed.

City - Which city is the insured.

Region - The region in which the city falls under.

```
all_health_df.isnull().sum()
```

```
age      0
sex      0
bmi      0
children 0
smoker   0
charges  0
city     0
region   0
dtype: int64
```

```
all_health_df.describe()
```

	age	bmi	children	charges
count	1338.000000	1338.000000	1338.000000	1338.000000
mean	39.207025	30.663397	1.094918	13270.422265
std	14.049960	6.098187	1.205493	12110.011237
min	18.000000	15.960000	0.000000	1121.873900
25%	27.000000	26.296250	0.000000	4740.287150
50%	39.000000	30.400000	1.000000	9382.033000
75%	51.000000	34.693750	2.000000	16639.912515
max	64.000000	53.130000	5.000000	63770.428010

```
print("**Unique Values By Features**")
for feature in all_health_df.columns:
    uniq = np.unique(all_health_df[feature])
    print(f" {feature} = {len(uniq)}")
```

```
**Unique Values By Features**
feature = 'age', 47
feature = 'sex', 2
feature = 'bmi', 548
feature = 'children', 6
feature = 'smoker', 2
feature = 'charges', 1337
feature = 'city', 1214
feature = 'region', 4
```

Summary of Categorical Features

```
all_health_df.describe(include=['object']).T
```

	count	unique	top	freq
sex	1338	2	male	676
smoker	1338	2	no	1064

city	1338	1214	Greenfield	4
region	1338	4	northeast	440

Replace missing values with defaults based on column

```
#only fill if there is any null values
if all_health_df.isna().sum().values.sum() > 0 :
    all_health_df['age'].fillna(all_health_df['age'].mean(),
inplace=True)
    all_health_df['smoker'].fillna("no", inplace=True)
    all_health_df['children'].fillna(0, inplace=True)
    all_health_df['sex'].fillna("male", inplace=True)


#below columns cannot be negative
all_health_df[all_health_df['age'] < 0]['age'].value_counts()
all_health_df[all_health_df['children'] < 0]
['children'].value_counts()
all_health_df[all_health_df['charges'] < 0]['charges'].value_counts()

Series([], Name: count, dtype: int64)
```

Include a histogram of each of the 5 variables – in your summary and analysis, identify any outliers and explain the reasoning for them being outliers and how you believe they should be handled (Chapter 2).


```
def create_histogram(df, label_1,label_2='frequency'):
    thinkplot.Hist(thinkstats2.Hist(df, label=label_1))
    thinkplot.show( xlabel=label_1, ylabel=label_2)

#Children
create_histogram(all_health_df['children'], "children", "count");
```




<Figure size 800x600 with 0 Axes>

```
#bmi  
create_histogram(all_health_df['bmi'], "bmi");
```




<Figure size 800x600 with 0 Axes>

```
#region  
create_histogram(all_health_df['region'], "region", "count");
```




<Figure size 800x600 with 0 Axes>

```
#smoker  
create_histogram(all_health_df['smoker'], "Smoker");
```



<Figure size 800x600 with 0 Axes>

```
create_histogram(all_health_df['sex'], "sex");
```



<Figure size 800x600 with 0 Axes>

Identify Outliers via Standard Deviation method

```
from collections import Counter

def StDev_method (df,n,features):
    """
    Takes a dataframe df of features and returns an index list
    corresponding to the observations
    containing more than n outliers according to the standard
    deviation method.
    """
    outlier_indices = []

    for column in features:
        # calculate the mean and standard deviation of the data frame
        feature_mean = df[column].mean()
        feature_std = df[column].std()

        # calculate the cutoff value, if the value is outside 3
        standard deviation
        cut_off = feature_std * 3
```



```

    # Determining a list of indices of outliers for feature column

    outlier_index = df[(df[column] < feature_mean - cut_off) |
(df[column] > feature_mean + cut_off)].index
    if len(outlier_index) > 0:
        print(f'{column=},{outlier_index=},{len(outlier_index)}')
        print(all_health_df.iloc[outlier_index][column])

    # appending the found outlier indices for column to the list
of outlier indices
    outlier_indices.extend(outlier_index)

    # selecting observations containing more than x outliers
    outlier_indices = Counter(outlier_indices)

    multiple_outliers = list( k for k, v in outlier_indices.items() if
v > n )

    # Calculate the number of records below and above lower and above
bound value respectively
    df1 = df[df[column] > feature_mean + cut_off]
    df2 = df[df[column] < feature_mean - cut_off]
    print('Total number of rows with outliers :', df1.shape[0]+
df2.shape[0])
    print(f'Total number of rows with more than {n} outliers =
{len(multiple_outliers)}, so we can safely say the dataset does not
have any major outliers.')
    return multiple_outliers

feature_list = ['age','bmi','charges','children']
# detecting outliers
max_outlier_per_row = 2
Outliers_StDev =
StDev_method(all_health_df,max_outlier_per_row,feature_list)

column='bmi',outlier_index=Index([116, 847, 1047, 1317],
dtype='int64'),4
116      49.06
847      50.38
1047     52.58
1317     53.13
Name: bmi, dtype: float64
column='charges',outlier_index=Index([34, 543, 577, 819, 1146, 1230,
1300], dtype='int64'),7
34      51194.55914
543      63770.42801
577      58571.07448
819      55135.40209
1146     52590.82939

```

```

1230      60021.39897
1300      62592.87309
Name: charges, dtype: float64
column='children',outlier_index=Index([ 32,  71, 166, 413, 425,
438, 568, 640, 877, 932, 937, 969,
      984, 1085, 1116, 1130, 1245, 1272],
      dtype='int64'),18
32         5
71         5
166        5
413        5
425        5
438        5
568        5
640        5
877        5
932        5
937        5
969        5
984        5
1085       5
1116       5
1130       5
1245       5
1272       5
Name: children, dtype: int64
Total number of rows with outliers : 18
Total number of rows with more than 2 outliers = 0, so we can safely
say the dataset does not have any major outliers.

```

Include the other descriptive characteristics about the variables: Mean, Mode, Spread, and Tails (Chapter 2). Using pg. 29 of your text as an example, compare two scenarios in your data using a PMF. Reminder, this isn't comparing two variables against each other – it is the same variable, but a different scenario. Almost like a filter. The example in the book is first babies compared to all other babies, it is still the same variable, but breaking the data out based on criteria we are exploring (Chapter 3).

```

def smoker_pmf(smkr, non_smkr):
    """Create PMF to compare
    """

    smkr_pmf = thinkstats2.Pmf(smkr.charges, label='smoker charges')
    non_smkr_pmf = thinkstats2.Pmf(non_smkr.charges, label='non smoker
charges')
    width = 0.45

    thinkplot.PrePlot(2, cols=2)
    thinkplot.Hist(smkr_pmf, align='right', width=width)
    thinkplot.Hist(non_smkr_pmf, align='left', width=width)

```

```


thinkplot.show()


thinkplot.PrePlot(2)
thinkplot.SubPlot(2)
thinkplot.Pmfs([smkr_pmf, non_smkr_pmf])
thinkplot.show()

smoker_data = all_health_df[all_health_df['smoker'] == 'yes']
non_smoker_data = all_health_df[all_health_df['smoker'] == 'no']

smoker_pmf(smoker_data, non_smoker_data);

```






<Figure size 800x600 with 0 Axes>

Create 1 CDF with one of your variables, using page 41-44 as your guide, what does this tell you about your variable and how does it address the question you are trying to answer (Chapter 4).

```
cdf = thinkstats2.Cdf(all_health_df.charges, label='Charges')
thinkplot.Cdf(cdf)
thinkplot.Config(xlabel='Charges', ylabel='CDF', loc='lower right');
```



Plot 1 analytical distribution and provide your analysis on how it applies to the dataset you have chosen (Chapter 5).

*#create dummies of categorical features so that correlation may be analyzed*

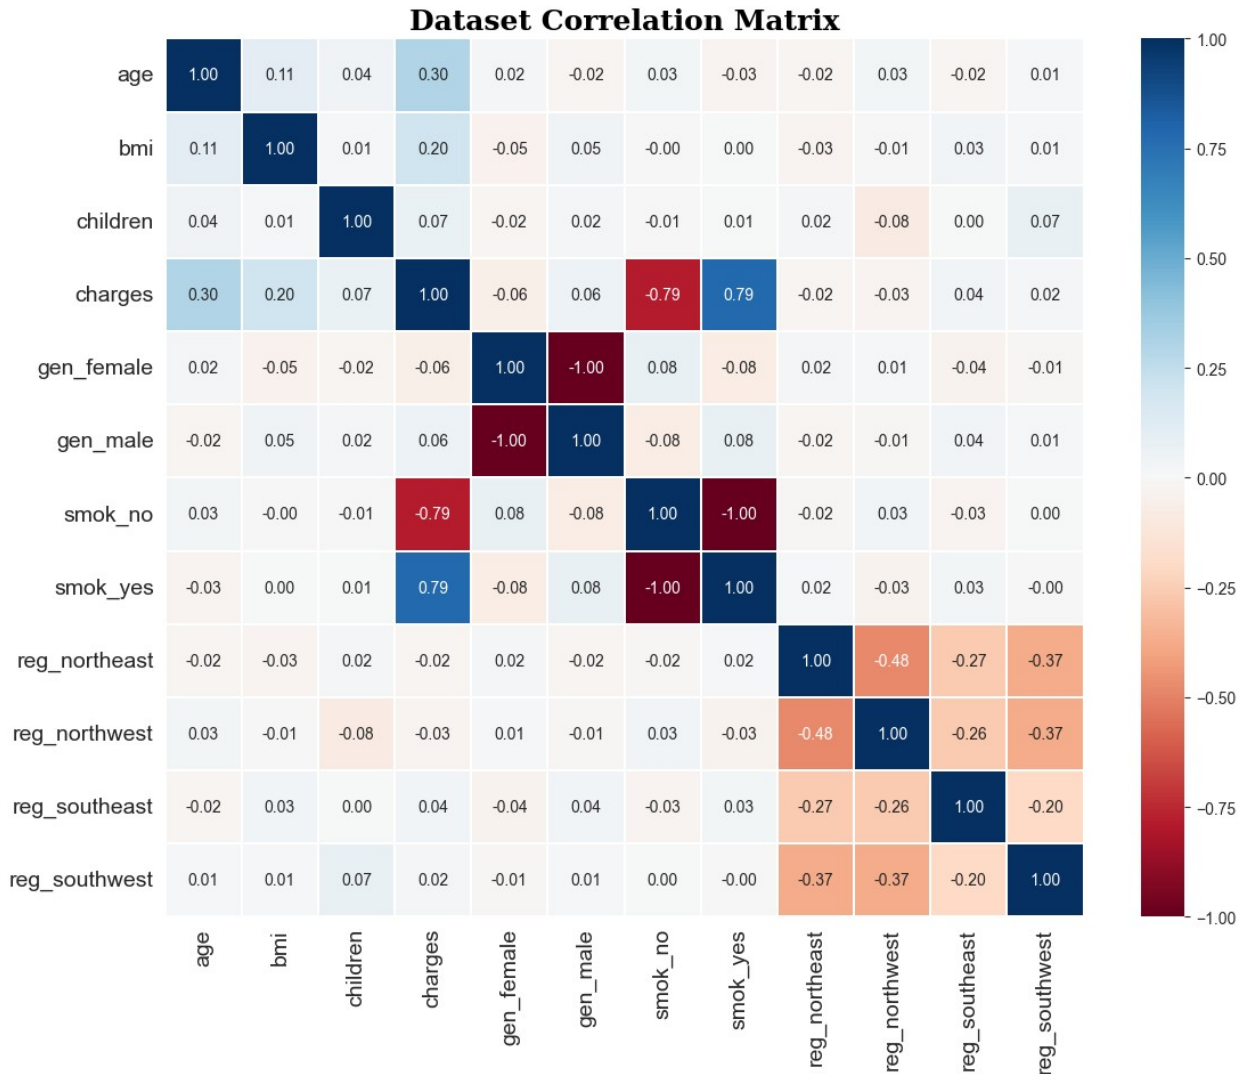
```
category_df= pd.get_dummies(all_health_df, prefix='gen',
columns=['sex'])
category_df= pd.get_dummies(category_df, prefix='smok',
columns=['smoker'])
category_df= pd.get_dummies(category_df, prefix='reg',
columns=['region'])
category_df.drop(columns='city',inplace=True)
category_df.head(10)
```

	age	bmi	children	charges	gen_female	gen_male	
smok_no \							
0	19	27.900	0	16884.92400	True	False	False
1	18	33.770	1	1725.55230	False	True	True
2	28	33.000	3	4449.46200	False	True	True
3	33	22.705	0	21984.47061	False	True	True

4	32	28.880	0	3866.85520	False	True	True
5	31	25.740	0	3756.62160	True	False	True
6	46	33.440	1	8240.58960	True	False	True
7	37	27.740	3	7281.50560	True	False	True
8	37	29.830	2	6406.41070	False	True	True
9	60	25.840	0	28923.13692	True	False	True

	smok_yes	reg_northeast	reg_northwest	reg_southeast	reg_southwest
0	True	True	False	False	False
1	False	False	True	False	False
2	False	False	True	False	False
3	False	False	False	True	False
4	False	False	False	False	True
5	False	False	True	False	False
6	False	False	True	False	False
7	False	True	False	False	False
8	False	False	False	True	False
9	False	True	False	False	False

```
#plot correlation matrix heatmap
fig, ax = plt.subplots(figsize=[13, 10])
sns.heatmap(category_df.corr(), ax=ax, annot=True, linewidths=0.05,
fmt='.2f', cmap="RdBu")
ax.tick_params(axis='both', which='major', labelsize=14)
ax.set_title('Dataset Correlation Matrix', fontdict={'family':
'serif', 'color': 'black', 'size': 18, 'weight': 'bold'});
```



```
#find which features have significant correlation (-0.5 < coef. > 0.5)

req_category_df = np.logical_or(category_df.corr() > 0.5,
category_df.corr() < -0.5)

req_category_df.head()
#disregard identity
for i in np.arange(0, len(req_category_df)):
    req_category_df.iloc[i,i] = False

#create table of correlation relationships by index values
corr_val = []
for i in np.arange(0, len(req_category_df.iloc[0, :])):
    for j in np.arange(0, len(req_category_df.iloc[:, 0])):
        if req_category_df.iloc[i,j] == True:
            corr_val.append([req_category_df.index.values[j],
req_category_df.columns.values[i]])
```

```

#drop rows in table in which relationships are duplicated in table (x
vs. y and y vs. x)
x = []
for i in np.arange(0,len(corr_val)):
    x.append(str(i))
corr_val = pd.DataFrame(corr_val, columns=["", ""], index=x)
del req_category_df, x
for i in corr_val.index:
    for j in corr_val.index:
        if any(corr_val.index == i) == True:
            a, b = corr_val.loc[i]
            if i != j:
                if np.logical_and((corr_val.loc[j][0] == b) == True,
(corr_val.loc[j][1] == a) == True) == True:
                    corr_val.drop(j, inplace=True);
corr_val.reset_index(drop=True, inplace=True)

print('The following features in the data are significantly correlated
(-0.5 < coef. > 0.5):\n {}'.format(corr_val))


The following features in the data are significantly correlated (-0.5
< coef. > 0.5):

0    smok_no    charges
1    smok_yes    charges
2    gen_male    gen_female
3    smok_yes    smok_no


sns.pairplot(all_health_df, corner=True);

```

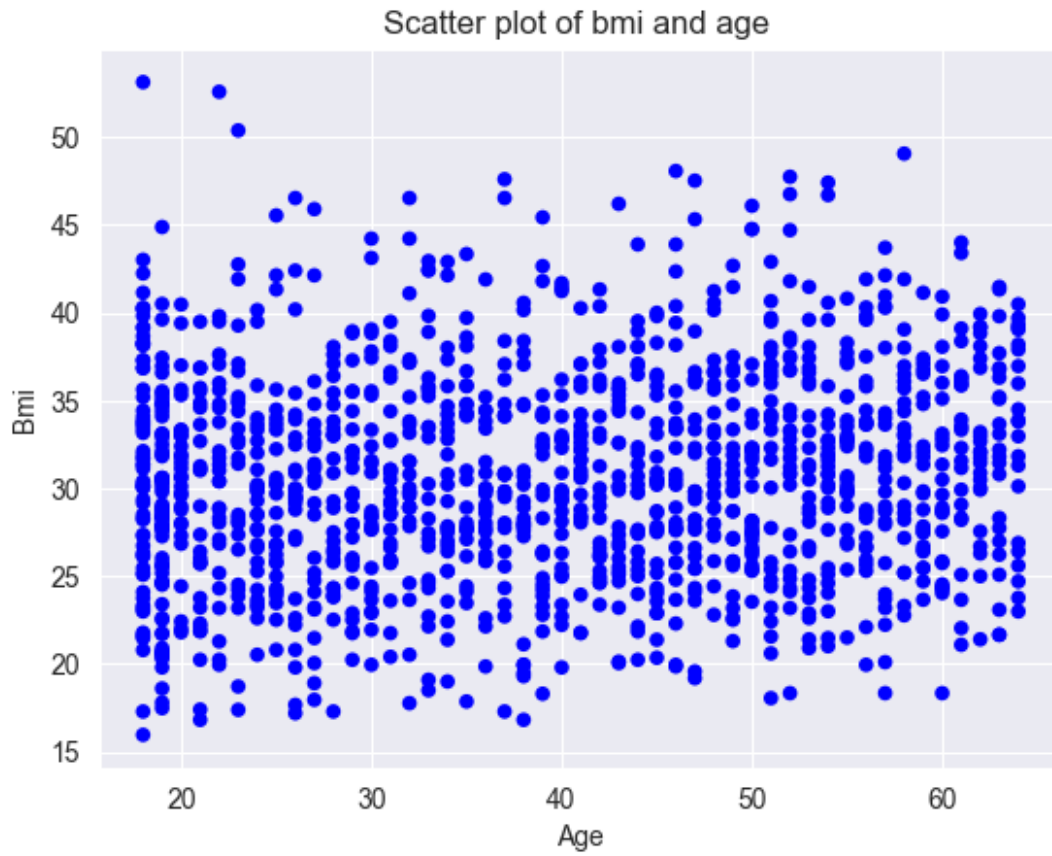




```
thinkplot.Scatter(all_health_df['age'], all_health_df['charges'],
alpha = 1)
thinkplot.Config(title = 'Scatter plot of Charges and age ',
xlabel = 'Age',
ylabel = 'Charges',
legend = False);
```



```
thinkplot.Scatter(all_health_df['age'], all_health_df['bmi'], alpha =  
1)  
thinkplot.Config(title = 'Scatter plot of bmi and age ',  
xlabel = 'Age',  
ylabel = 'Bmi',  
legend = False);
```



Conduct a test on your hypothesis using one of the methods covered in Chapter 9.

## Hypothesis Testing

```
def run_hypothesis_test(data, iters=1000):
    """Test the difference in means between Alaska Airlines and other
    airlines
    """

    # test the difference in means
    ht = hypothesis.DiffMeansPermute(data)
    p_value = ht.PValue(iters=iters)
    print('\nmeans permute two-sided')
    hypothesis.PrintTest(p_value, ht)

    ht.PlotCdf()

    # test the difference in means one-sided
    ht = hypothesis.DiffMeansOneSided(data)
    p_value = ht.PValue(iters=iters)
    print('\nmeans permute one-sided')
```

```
hypothesis.PrintTest(p_value, ht)
```

```
# test the difference in std
```

```
ht = hypothesis.DiffStdPermute(data)
```

```
p_value = ht.PValue(iters=iters)
```

```
print('\nstd permute one-sided')
```

```
hypothesis.PrintTest(p_value, ht)
```

```
hypothesisTestData= smoker_data.charges.values,  
non_smoker_data.charges.values  
run_hypothesis_test(hypothesisTestData)
```

```
0 1000
```

```
means permute two-sided
```

```
p-value = 0.0
```

```
actual = 23615.963533676637
```

```
ts max = 3263.261602219054
```

```
0 1000
```

```
means permute one-sided
```

```
p-value = 0.0
```

```
actual = 23615.963533676637
```

```
ts max = 2895.1017710675733
```


```
0 1000
```

```
std permute one-sided
```

```
p-value = 0.0
```

```
actual = 5529.502176241819
```

```
ts max = 2524.594327813067
```



Conclusion: The p-value is near to 0 , which means it is plausible that the observed difference is just the result of random sampling, and might not be generally true in the population. Reject the Null Hypothesis.

## Hypothesis Testing

```
hypothesisTestData= smoker_data.bmi.values, non_smoker_data.bmi.values
```

```
run_hypothesis_test(hypothesisTestData)
```

```
891 1000
```

```
means permute two-sided
```

```
p-value = 0.891
```

```
actual = 0.05665379232753409
```

```
ts max = 1.4330402420284258
```

```
456 1000
```

```
means permute one-sided
```


```
p-value = 0.456
```

```
actual = 0.05665379232753409
```

```
ts max = 1.3021251920860522
```

```
172 1000
```

```
std permute one-sided
p-value = 0.172
actual = 0.26683215652024117
ts max = 1.0362670517620627
```



p-value of .44 is significant than alpha of .05 and hence Null hypothesis cannot be rejected. In other words, the claims made by smokers is less than or equal to that of non-smokers.

```
def linear_reg_test(all_data,x_data,y_data):
    """Plots a scatter plot and fitted curve.
    """
    sample = thinkstats2.SampleRows(all_data, 1000)
    inter, slope = thinkstats2.LeastSquares(x_data, y_data)
    fit_xs, fit_ys = thinkstats2.FitLine(x_data, inter, slope)

    thinkplot.Scatter(x_data, y_data, color='gray', alpha=0.1)
    thinkplot.Plot(fit_xs, fit_ys, color='white', linewidth=3)
    thinkplot.Plot(fit_xs, fit_ys, color='blue', linewidth=2)
    thinkplot.show(root='chargesbmi linear',
                    xlabel='bmi',
                    ylabel='charges',
                    # axis=[10, 45, 0, 15],
                    legend=False)


formula = 'charges ~ bmi'
```

```

model = smf.ols(formula, data=sample)
results = model.fit()
regression.SummarizeResults(results)

bmi_df,charges_df = all_health_df['bmi'],all_health_df['charges']
linear_reg_test(all_health_df,bmi_df,charges_df)

```



```

Intercept    -568    (0.773)
bmi          458    (6.72e-13)
R^2  0.05044
Std(ys) 1.235e+04
Std(res) 1.204e+04

<Figure size 800x600 with 0 Axes>

formula = 'charges ~ bmi + children + age'
model = smf.ols(formula, data = all_health_df)
results = model.fit()
results.summary()

<class 'statsmodels.iolib.summary.Summary'>
"""

```

OLS Regression Results

```

=====
=====
Dep. Variable:          charges    R-squared:
0.120
Model:                  OLS      Adj. R-squared:
0.118
Method:                 Least Squares    F-statistic:
60.69
Date:                   Sat, 11 Jan 2025    Prob (F-statistic):
8.80e-37
Time:                   19:29:29    Log-Likelihood:
-14392.
No. Observations:       1338    AIC:
2.879e+04
Df Residuals:           1334    BIC:
2.881e+04
Df Model:                3

```

Covariance Type: nonrobust

```

=====
=====
              coef      std err          t      P>|t|      [0.025
0.975]
-----
-----
Intercept  -6916.2433    1757.480     -3.935     0.000    -1.04e+04    -
3468.518
bmi         332.0834     51.310      6.472     0.000     231.425
432.741
children    542.8647     258.241      2.102     0.036      36.261
1049.468
age         239.9945     22.289     10.767     0.000     196.269
283.720

```

```

=====
=====
Omnibus:              325.395    Durbin-Watson:
2.012
Prob(Omnibus):        0.000    Jarque-Bera (JB):
603.372
Skew:                 1.520    Prob(JB):
9.54e-132
Kurtosis:              4.255    Cond. No.
290.

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is



```
correctly specified.  
"""
```